

# Exploratory Data Analysis

## Introduction

---

The primary goal of **Exploratory Data Analysis (EDA)** is to gain a comprehensive understanding of the data, its quality, identify any underlying trends or relationships that may influence the analysis and modeling process.

### Importance:

- EDA helps detect outliers, and inconsistencies in the data, addressing these issues ensures that the analysis is based on accurate and clean data.
- Through visualizations and statistical analysis, EDA identifies which features are most relevant to the questions and should be included in the model.
- EDA reveals relationships between variables, such as correlations, and trends, which can help discard certain features or create new derived features.
- By analyzing the distribution of variables, we can identify central tendencies and variations. This helps us understand how the data is spread out and if it meets assumptions required by certain algorithms.

### Techniques:

- Univariate Analysis:
  - To understand the distribution, and central tendency of a variable.
  - Visualizations: Histograms, Box Plots, kernel Density plots.
  - Statistical measures: Mean, Median, Mode, Variance, Standard Deviation.
- Bivariate Analysis:
  - To understand the relationship between two variables
  - Categorical-Categorical: Heatmaps.
  - Numerical-Categorical: Box plots, Bar Graphs.
  - Numerical-Numerical: Scatter plots, correlation coefficients.
- Multivariate Analysis:
  - To explore relationships between three or more variables simultaneously.
  - Visualizations: Heatmaps, Scatter plots.
  - Dimensionality Reduction: PCA, t-SNE.
- Data Distribution:
  - Histograms and density plots are useful to detect skewness or multimodal distributions.
- Statistical Analysis:

- Hypothesis Testing: t-tests, chi-square tests, ANOVA.
- Correlation: Pearson, Spearman correlation coefficients
- Summary Statistics: Mean, variance, Quartiles.

### In this project:

There are various datasets, and multiple independent features such as: Driver Performance, Pit stop durations, race track features, overall performance of the car, which play an important role in determining the outcome of the race. Therefore, EDA is crucial to understand how these factors influence race results.

## Code

Provide the source code used for this section of the project here.

If you're using a package for code organization, you can import it at this point. However, make sure that the **actual workflow steps**—including data processing, analysis, and other key tasks—are conducted and clearly demonstrated on this page. The goal is to show the technical flow of your project, highlighting how the code is executed to achieve your results.

If relevant, link to additional documentation or external references that explain any complex components. This section should give readers a clear view of how the project is implemented from a technical perspective.

Remember, this page is a technical narrative, NOT just a notebook with a collection of code cells, include in-line Prose, to describe what is going on.

## Required Libraries

---

```
# import required libraries
from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt
import json
import os
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.lines as mlines
from matplotlib.cm import get_cmap
import seaborn as sns
from scipy.stats import ttest_ind
```

## News Data<sup>1</sup>

---

In this section, we will analyze media coverage of select F1 drivers who created significant buzz during the 2024 season. The drivers include **Max Verstappen, Carlos Sainz, Lando Norris, Daniel Riccardo**, and **Lewis Hamilton**. These drivers have been at the center of media discussions due to their performance, contract negotiations, and potential moves for the 2025 season.

```

# Function to generate a word cloud for a single JSON file
def generate_wordcloud_top_n(file_path, output_folder, top_n=10):
    # helper function to plot and save the word cloud
    def plot_cloud(wordcloud, output_file):
        plt.figure(figsize=(10, 8))
        plt.imshow(wordcloud, interpolation='bilinear')
        plt.axis("off")
        plt.savefig(output_file)

        plt.close()

    # read the text from the JSON file
    with open(file_path, 'r', encoding='utf-8') as file:
        data = json.load(file)

    # concatenate all text content if it's a JSON dictionary
    if isinstance(data, dict):
        my_text = ' '.join(data.values())

    # generate word cloud to extract word frequencies
    wordcloud = WordCloud(
        width=3000,
        height=2000,
        random_state=1,
        background_color="white",
        colormap="Blues_r",
        collocations=False,
        stopwords=STOPWORDS
    ).generate(my_text)

    # extract word frequencies and get the top N words
    word_frequencies = wordcloud.words_
    top_words = dict(list(word_frequencies.items())[:top_n])

    # generate a new word cloud with only the top N words
    top_wordcloud = WordCloud(
        width=3000,
        height=2000,
        random_state=1,
        background_color="white",
        colormap="Blues_r",
        collocations=False,
        stopwords=STOPWORDS
    ).generate_from_frequencies(top_words)

    # Save and display the word cloud
    file_name = os.path.basename(file_path).replace('.json', '_wordcloud.png')
    output_file = os.path.join(output_folder, file_name)
    plot_cloud(top_wordcloud, output_file)

# function to process all JSON files in a folder
def generate_wordclouds_for_folder(input_folder, output_folder, top_n=10):
    os.makedirs(output_folder, exist_ok=True)
    # loop through all files in the input folder
    for file_name in os.listdir(input_folder):

```

```
# process only JSON files
if file_name.endswith('.json'):
    file_path = os.path.join(input_folder, file_name)
    try:
        generate_wordcloud_top_n(file_path, output_folder, top_n=20)
    # for debugging purpose
    except Exception as e:
        print(f"Error processing file {file_path}: {e}")
```

```
input_folder = "../../../data/processed-data/News_Drivers/"
output_folder = "../../../data/eda/WordClouds/"
generate_wordclouds_for_folder(input_folder, output_folder, top_n=20)
```

## Analysis

### Max Verstappen



- Strong occurrence of “Norris” suggests the media discussion about the rivalry and competition between the top two drivers.
- The presence of “penalty” indicates that articles may have covered controversies or penalties during races that involved Verstappen. This adds to the narrative of the challenges and incidents faced during his title run.
- Words such as “title” and “world champion” reinforce the focus on Max Verstappen’s achievements, particularly his dominance throughout the season and his consecutive championship wins.

## Drivers’ Performance

```
# number of points by driver
data = pd.read_csv("../data/processed-data/driver_standings_2000_2020.csv")
# Calculate total points for each driver
total_points = data.groupby("driverName")["Points"].sum().reset_index()

# plot only top 10 drivers
top_10_drivers = total_points.sort_values(by="Points", ascending=False).head(10)

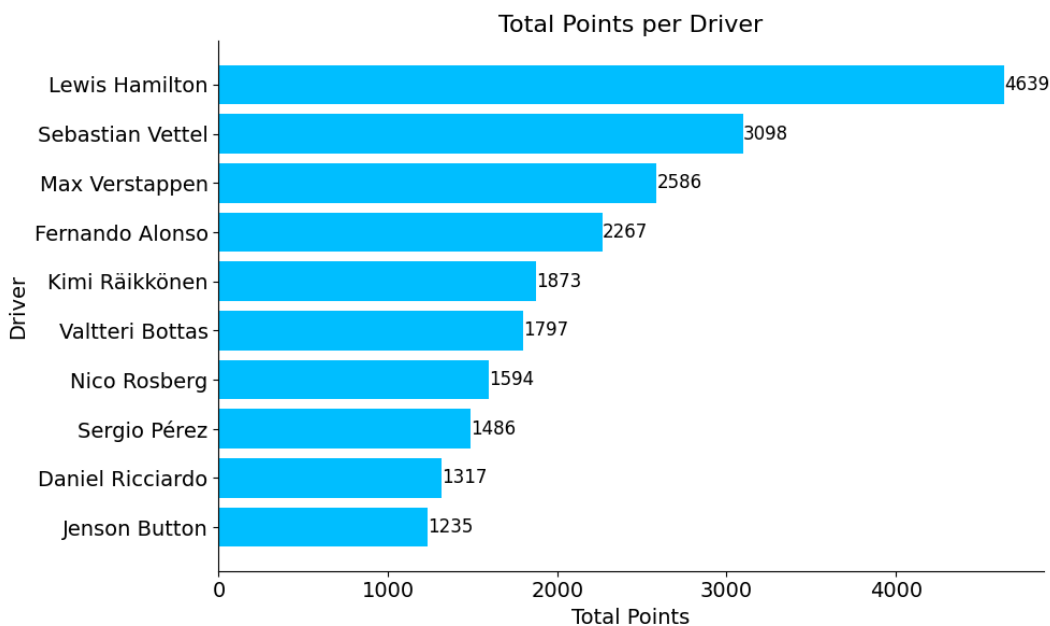
# Plot the bar chart
plt.figure(figsize=(10, 6))
bars = plt.barh(top_10_drivers["driverName"], top_10_drivers["Points"])

for bar in bars:
    plt.text(
        bar.get_width() + 1,
        bar.get_y() + bar.get_height() / 2,
        f'{int(bar.get_width())}',
        va='center', fontsize=12, color='black'
    )
```

```
plt.xticks(fontsize = 14)
plt.yticks(fontsize = 14)
plt.xlabel("Total Points", fontsize=14)
plt.ylabel("Driver", fontsize=14)
plt.title("Total Points per Driver", fontsize=16)
plt.gca().invert_yaxis()
plt.tight_layout()

plt.gca().spines['top'].set_visible(False)
plt.gca().spines['right'].set_visible(False)

plt.show()
```



#### Insights:

- The cumulative points reflect a driver's career longevity, performance, and overall success in the sport
- **Lewis Hamilton** stands out as the clear leader with 4,630 points, significantly ahead of other drivers. This highlights his dominance in the sport over the years, attributed to his consistent race wins, and career with top teams like McLaren and Mercedes in their eras.
- **Sebastian Vettel** holds the second spot with 3,098 points, his performance in Red Bull and Ferrari might have significant contribution to these numbers, since his early years in Sauber, and Toro Rosso, and his performance in Aston Martin were not as successful.
- **Max Verstappen** has the third highest cumulative total despite entering the sport relatively later. His rapid accumulation of points in the recent years underscores his dominance in the current era of F1.

```
# points scored by driver over time with different teams
def plot_driver_points(data, driver_name, color_list=None):

    # filter the data to include only the specified driver
    # sort by season
    driver_data = data[data['driverName'] == driver_name].sort_values('season')
    # list of unique constructors the driver raced for
```

```

unique_constructors = driver_data["Constructor_Name"].unique()
# assign colors to constructors
if color_list:
    constructor_colors = {
        constructor: color_list[i % len(color_list)]
        for i, constructor in enumerate(unique_constructors)
    }
else:
    constructor_colors = {
        constructor: plt.cm.tab10(i / len(unique_constructors))
        for i, constructor in enumerate(unique_constructors)
    }
plt.figure(figsize=(10, 6))

# line plot
for i in range(len(driver_data) - 1):
    season_start = driver_data.iloc[i]['Season']
    season_end = driver_data.iloc[i + 1]['Season']
    points_start = driver_data.iloc[i]['Points']
    points_end = driver_data.iloc[i + 1]['Points']
    constructor = driver_data.iloc[i]['Constructor_Name']

    plt.plot(
        [season_start, season_end],
        [points_start, points_end],
        color=constructor_colors[constructor],
        linewidth=4
    )

for _, row in driver_data.iterrows():
    plt.text(row['Season'], row['Points'], str(row['Points']),

plt.title(f"Points Over Seasons for {driver_name}", fontsize=16)
plt.xticks(fontsize = 14, rotation = 90)
plt.yticks(fontsize = 14)
plt.xlabel("Season", fontsize=14)
plt.ylabel("Points", fontsize=14)
plt.tight_layout()
plt.xticks(range(int(driver_data["Season"].min()), int(driver_data["Season"].max()) + 1),

legend_elements = [
    mlines.Line2D([], [], color=constructor_colors[constructor],
    for constructor in unique_constructors
]
plt.legend(handles=legend_elements, loc="upper left", fontsize=12)

plt.show()




```




```

color_list = ["lightskyblue", "cornflowerblue", "deepskyblue", "powderblue"]

```

```
plot_driver_points(data, "Lewis Hamilton", color_list)
```





Insights:

- A significant portion of the total championship points for most drivers has been earned while driving for top-performing teams such as Ferrari and Red Bull.



- There is a clear correlation between the points scored in a season and the team a driver represented, emphasizing the importance of team performance in a driver's success.
- Observing the career trajectories of drivers like Daniel Riccardo, Kimi Räikkönen, and Sebastian Vettel, we notice a pattern where they initially drove for top teams, followed by a move to midfield teams before retiring.
- Some drivers switched teams despite scoring a substantial number of championship points in the previous season. This highlights the strategic decisions made by teams in planning their driver line-ups and the drivers' foresights regarding the potential improvements and competitiveness of cars across the grid.

## Constructors' Performance

```
total_points = data.groupby("Constructor_Name")["Points"].sum().reset_index()


constructors = total_points.sort_values(by="Points", ascending=False)

# Plot the bar chart
plt.figure(figsize=(10, 6))
bars = plt.barh(constructors["Constructor_Name"], constructors["Points"])

for bar in bars:
    plt.text(
        bar.get_width() + 1,
        bar.get_y() + bar.get_height() / 2,
        f'{int(bar.get_width())}',
        va='center', fontsize=12, color='black'
    )
plt.xticks(fontsize = 14)
plt.yticks(fontsize = 14)
plt.xlabel("Total Points", fontsize=14)
plt.ylabel("Constructor", fontsize=14)
plt.title("Total Points per Constructor", fontsize=16)
plt.gca().invert_yaxis()
plt.tight_layout()

plt.gca().spines['top'].set_visible(False)
plt.gca().spines['right'].set_visible(False)

plt.show()
```



#### Insights:

- The top 3 teams – Ferrari, Mercedes, and Red Bull – are closely positioned in the cumulative points standings.
- Mercedes' position is particularly notable because they re-entered the sports as a constructor in 2010, while Ferrari and Red Bull have accumulated points across the full 2000-2023 period. Despite the shorter timeframe, Mercedes managed to secure the second position, showcasing their unparalleled dominance 2014 onward.
- From domain knowledge, a key differentiating factor for Mercedes has been the consistent performance of their drivers, such as Lewis Hamilton, Nico Rosberg, and Valtteri Bottas, who regularly secured podium finishes during their tenure. This level of consistency has allowed Mercedes to close the gap with Ferrari and Red Bull in a relatively short period.
- For Ferrari and Red Bull, the points reflect their long-standing success and ability to remain competitive across different eras. Ferrari's consistent presence at the top, coupled with Red Bull's dominance during Sebastian Vettel's era (2010-2013) and Max Verstappen's current reign, has solidified their positions in the top three.

```
# distribution of "Lapped", "Mechanical", "Accident" across Constructors
data = pd.read_csv("../data/processed-data/race_info.csv")
```

```
# Filter the data for specific statuses
filtered_data = data[data["status"].isin(["Lapped", "Mechanical", "Accident"])

# Group by constructor and status to get counts
status_distribution = filtered_data.groupby(["constructorName", "status"])

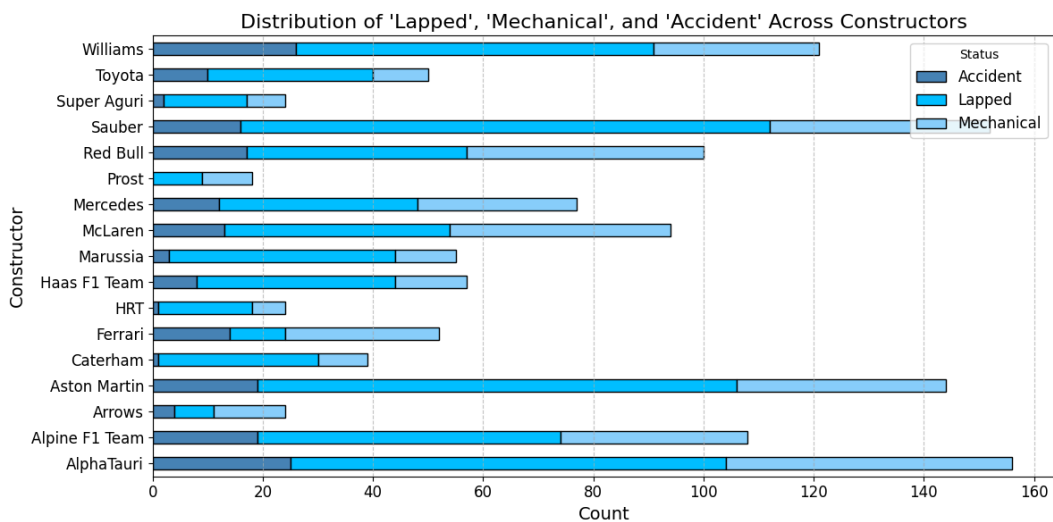
# Plot the distribution as a stacked horizontal bar chart
status_distribution.plot(
    kind="barh",
    figsize=(12, 6),
    color=['steelblue', 'deepskyblue', 'lightskyblue'],
```

```

    edgecolor="black",
    stacked=True
)

# Customize the plot
plt.title("Distribution of 'Lapped', 'Mechanical', and 'Accident' Across Constructors")
plt.xlabel("Count", fontsize=14)
plt.ylabel("Constructor", fontsize=14)
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)
plt.legend(title="Status", fontsize=12, loc="upper right")
plt.grid(axis="x", linestyle="--", alpha=0.7)
plt.tight_layout()
plt.show()

```



```

# distribution of status for each constructor
def autopct_format(pct):
    return f'{pct:.1f}%' if pct > 0 else ''

filtered_data = data[data["status"].isin(["Finished", "Lapped", "Mechanical"])
# Group by constructor and status to get counts
status_distribution = filtered_data.groupby(["constructorName", "status"])
# total number of entries per constructor
status_distribution["Total"] = data.groupby("constructorName").size()
# total number of constructors
num_constructors = len(status_distribution)
# for subplot
num_cols = 3
num_rows = (num_constructors + num_cols - 1) // num_cols
fig, axes = plt.subplots(num_rows, num_cols, figsize=(15, 5 * num_rows))

axes = axes.flatten()

colors = ['steelblue', 'deepskyblue', 'lightskyblue', 'powderblue']
# loop through each constructor and create a pie chart
for idx, constructor in enumerate(status_distribution.index):
    # extract status counts for each constructor
    status_counts = status_distribution.loc[constructor, ["Finished", "Lapped", "Mechanical"]]
    total_count = status_distribution.loc[constructor, "Total"]
    # calculate percentage for each status

```

```

percentages = (status_counts / total_count) * 100

wedges, texts, autotexts = axes[idx].pie(
    percentages,
    labels=None,
    autopct=autopct_format,
    colors=colors,
    startangle=140,
    textprops={'fontsize': 14}
)
axes[idx].set_title(f"{constructor}", fontsize=14)

for autotext in autotexts:
    autotext.set_fontsize(12)


for i in range(idx + 1, len(axes)):
    fig.delaxes(axes[i])

# Add a shared legend
fig.legend(
    labels=["Finished", "Lapped", "Mechanical", "Accident"],
    loc="lower center",
    ncol=1,
    fontsize=12
)


plt.tight_layout(rect=[0, 0.1, 1, 1])
plt.show()

```


AlphaTauri




Alpine F1 Team




Arrows




Aston Martin




Caterham




Ferrari




HRT




Haas F1 Team




Marussia




McLaren




Mercedes




Prost




Red Bull




Sauber




Super Aguri



Toyota




Williams





```
plt.title("Trends in Points for Selected Constructors", fontsize=16)
plt.xlabel("Season", fontsize=14)
plt.ylabel("Points", fontsize=14)
plt.xticks(grouped_data.index, rotation=90, fontsize=12)
plt.yticks(fontsize=12)
plt.legend(title="Constructor", fontsize=12)
plt.tight_layout()

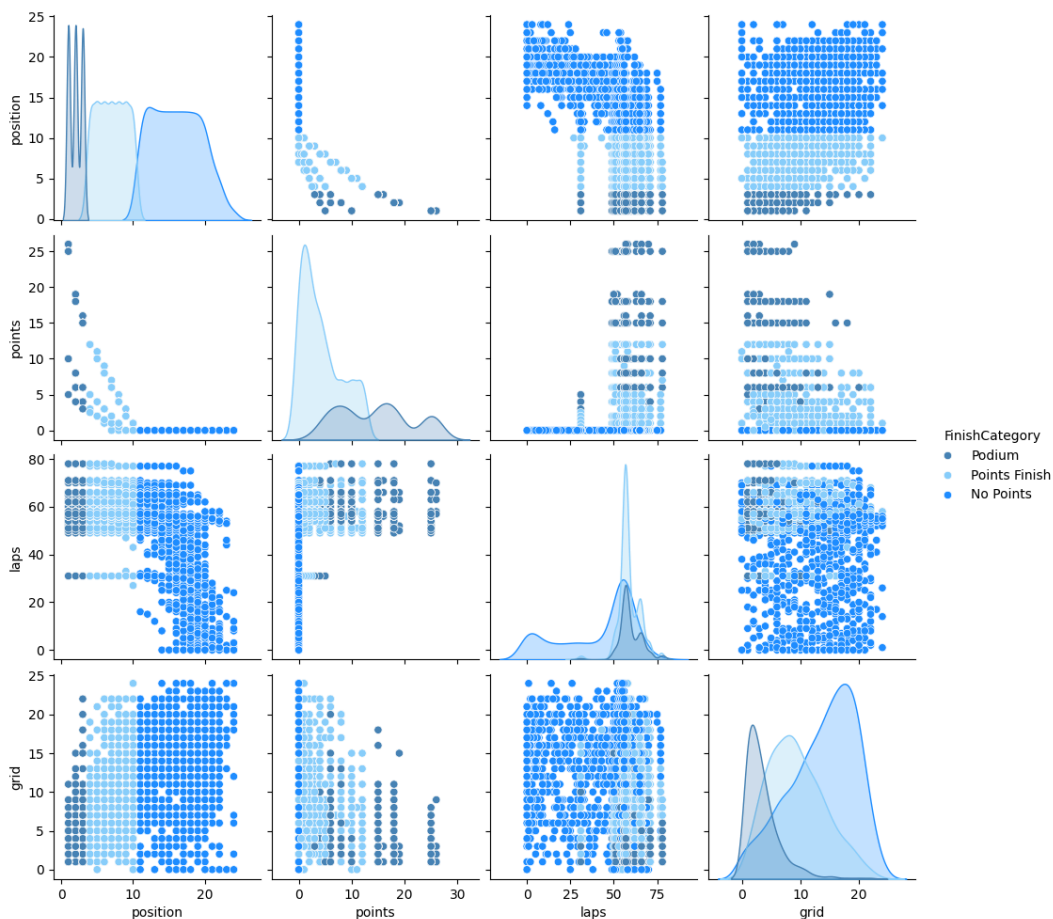
plt.show()
```



### Insights:

- The plot highlights clear periods of highs and lows for each team, emphasizing the dynamic nature of Formula 1. Success in the sport is cyclical, with teams experiencing periods of dominance followed by rebuilding phases.
- Red Bull showed a rapid rise starting in 2009, coinciding with their dominance during the Sebastian Vettel era (2010–2013). After a slight dip between 2014–2020, they have experienced another significant upward trend, driven by Max Verstappen’s dominance in recent years.
- Mercedes entered the sport as a constructor in 2010 and saw rapid progress starting in 2014 during the hybrid engine era, where they dominated for nearly a decade. Their decline starting in 2021 coincides with the resurgence of Red Bull and highlights challenges in adapting to regulation changes.
- Ferrari has experienced the most noticeable highs and lows among the teams. A steep decline is observed in 2020, likely due to car performance struggles, followed by a sharp recovery in 2022, marking their return to competitiveness.
- McLaren had strong results in the early 2000s but faced a significant decline after 2012, struggling during the hybrid era. Their gradual recovery post-2018, reflected in the upward trend, signals improvements in car performance and team structure.
- This analysis underscores that success in Formula 1 is not just about driver skill but also hinges on the ability of teams to innovate, adapt, and build competitive cars under constantly evolving technical regulations.

```
sns.pairplot(data[['position', 'points', 'laps', 'grid', 'FinishCategory']])
```



## Circuit Features

```
data = pd.read_csv("../data/processed-data/race_track_features.csv")
```

```
# Full Throttle (%) vs Number of Straights

# grouping based on the median of 'Number of Straights'
median_straights = data['Number of Straights'].median()

# tracks with fewer straights than median
low_straights = data[data['Number of Straights'] < median_straights]
# tracks with >= median straights
high_straights = data[data['Number of Straights'] >= median_straights]

# perform t-test
t_stat, p_value = ttest_ind(low_straights, high_straights, nan_policy='omit')

# results
print(f"T-Statistic: {t_stat:.4f}")
print(f"P-Value: {p_value:.4f}")
```

T-Statistic: 8.7778

P-Value: 0.0000



There is a significant difference in Full Throttle (%) between tracks with low and high straights.

```
# Full Throttle (%) vs Number of Corners

median_corners = data['Number of Corners'].median()

# tracks with lower number of corners than median
low_corners = data[data['Number of Corners'] < median_corners]['Full Throttle (%)']
# tracks with >= median corners
high_corners = data[data['Number of Corners'] >= median_corners]['Full Throttle (%)']

# perform t-test
t_stat, p_value = ttest_ind(low_corners, high_corners, nan_policy='omit')

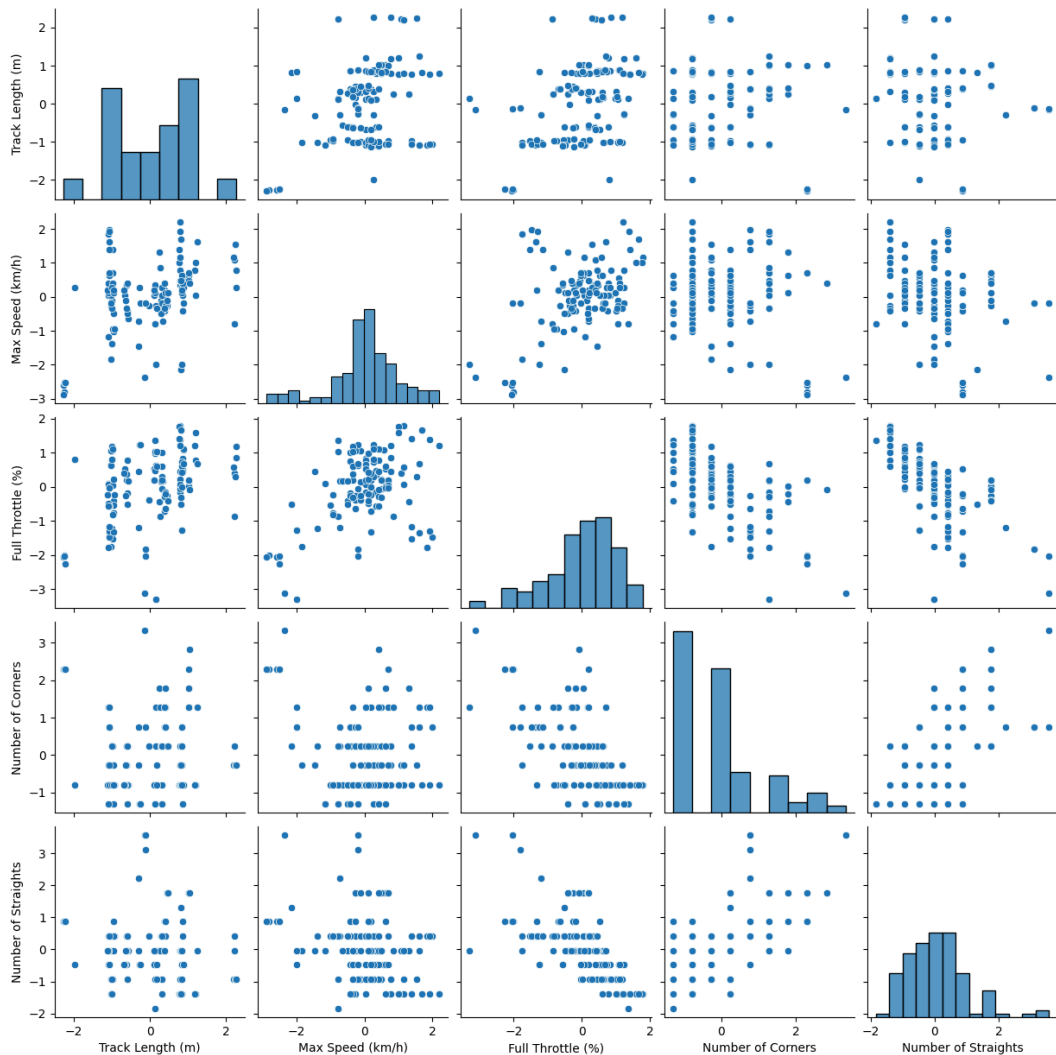
# results
print(f"T-Statistic: {t_stat:.4f}")
print(f"P-Value: {p_value:.4f}")
```

T-Statistic: 6.8204

P-Value: 0.0000

There is a significant difference in Full Throttle (%) between tracks with low and high corners.

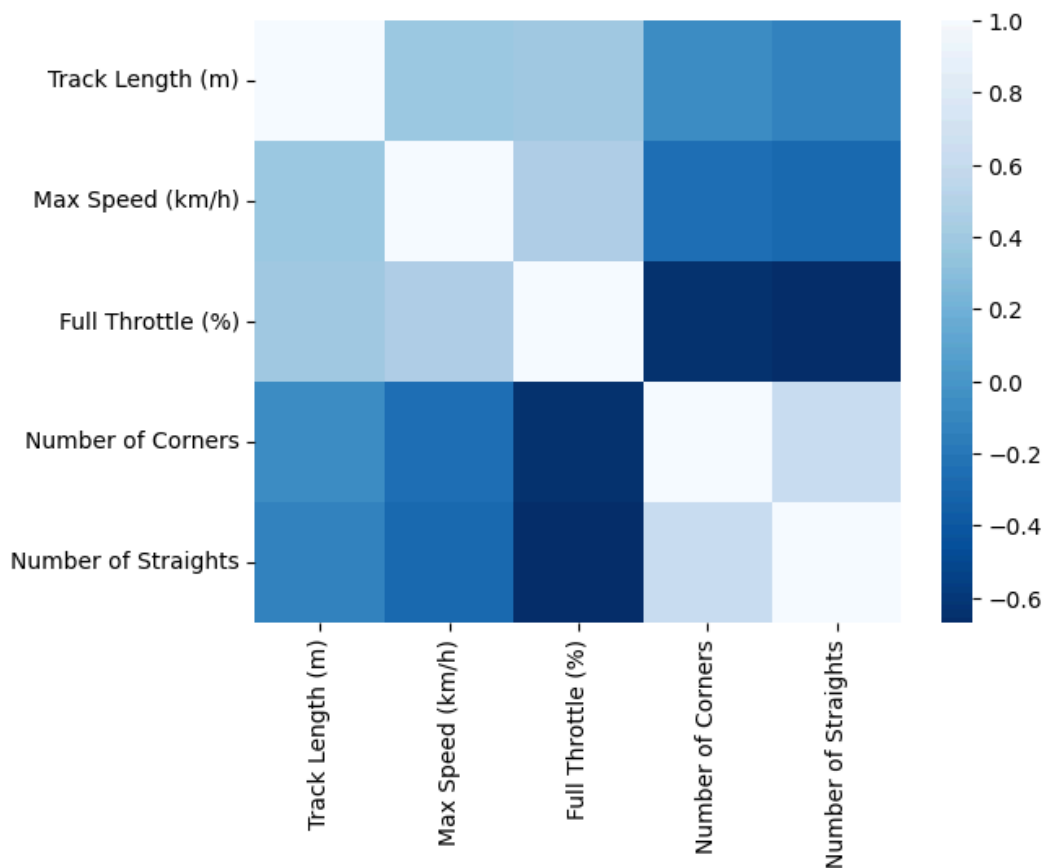
```
sns.pairplot(data[["Track Length (m)", "Max Speed (km/h)", "Full Throttle (%)"]])
plt.show()
```



```
corr_matrix = data[["Track Length (m)", "Max Speed (km/h)", "Full Throttle (%)", "Number of Corners", "Number of Straights")]
corr_matrix
```

	Track Length (m)	Max Speed (km/h)	Full Throttle (%)	Number of Corners	Number of Straights
Track Length (m)	1.000000	0.373616	0.399017	-0.058163	-0.118008
Max Speed (km/h)	0.373616	1.000000	0.462422	-0.246056	-0.281559
Full Throttle (%)	0.399017	0.462422	1.000000	-0.654084	-0.667532
Number of Corners	-0.058163	-0.246056	-0.654084	1.000000	0.631757
Number of Straights	-0.118008	-0.281559	-0.667532	0.631757	1.000000

```
sns.heatmap(corr_matrix, cmap="Blues_r")
plt.show()
```



## Pitstop Analysis

```
data = pd.read_csv("../data/processed-data/pitstop_with_positions.csv")
data.head()
```


Unnamed: 0									
0	Year	Round	RaceName	DriverID	Lap1	Lap2	Lap3	Lap4	Lap5
0 0	2011	1	Australian Grand Prix	alguersuari	-1.492522	-0.208201	1.339923	-0.208201	-0.208201
1 1	2011	1	Australian Grand Prix	alonso	-0.414215	0.355923	1.723412	-0.414215	-0.414215
2 2	2011	1	Australian Grand Prix	ambrosio	-0.218159	0.976458	-0.577525	-0.577525	-0.577525
3 3	2011	1	Australian Grand Prix	barrichello	-0.316187	0.130273	0.956433	0.956433	0.956433
4 4	2011	1	Australian Grand Prix	buemi	-0.120131	0.468747	-0.577525	-0.577525	-0.577525

5 rows × 36 columns

```
numerical_columns = ["Lap1", "Lap2", "Lap3", "Lap4", "Stop2", "Stop3", "Stop4", "Duration1", "Duration2", "Duration3", "Duration4"]
correlation_data = data[numerical_columns]

# Calculate correlation matrix
correlation_matrix = correlation_data.corr()
```

```
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap="Blues_r", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



Some of the cell outputs were excluded from the main page. They can be found [here](#)

### References

1. DSAN 5000 lecture content, <https://jfh.georgetown.domains/centralized-lecture-content/course-timelines/dsan-5000.html>.